

Task-1 Report

What I Learned From the Lab

The core lesson of the entire lab is one sentence: Get the right context in, keep the wrong context out. Every exercise was a variation on that rule.

- Context is everything. An AI is like a smart new coworker it only knows what you tell it. The difference between a generic answer and a specific is the amount of relevant context you provide it.
- Confidence is not correctness. AI can sound very certain even when guessing. For recent, local, or private information, I must ask it to search or admit uncertainty.
- You control the tool with your wording. Words like "today," "latest," and "cite sources" push towards web search. Words like "research thoroughly" push towards deep research. There is no mode button my wording picks the mode.
- Sycophancy is real. How you ask shapes the answer.
- The first answer is never the best. The Brainstorm Iterate loop is the highest-leverage prompting habit. Each iteration sharpens the output.
- Always force code for data work. Without specifying code, the AI may guess numbers. The fix: Write and run code, show me the code.
- Different models have different strengths. Keep two tabs open to uncover blind spots.

How My Prompting Improved During the Process

Before the lab: I was writing short, vague prompts and accepting whatever the AI gave back. I was not aware of how much the framing of my question affected answer quality.

After the lab: I now use structured brief prompts before I ask anything. I never accept the first answer, I iterate. I always check whether the AI ran code or just guessed numbers from text.

Prompting Techniques That Helped the Most

- Context is Everything (Exercise 4): "I need help with [task]. Here is what you need to know: [facts, constraints, audience]. What I want back: [format]." This single shape made every answer more useful from the first message.
- Stop the Flattery (Exercise 6): Instead of "Don't you think answer A is better?" Use "Compare answer A and B. List the strongest case for each. Eliminated the sycophancy problem.
- Brainstorm–Iterate loop (Exercise 7): Round 1: Ask for options. Round 2: Give feedback. Round 3: Expand the winner. The final output was always significantly better.

- Force code verification (Exercise 10): "Write and run code, show me the code you ran."
Turned guesses into verifiable results.
- Cross-model checking (Exercise 13): Running the same prompt on two different models revealed blind spots neither would have flagged alone.

Challenges I Faced and How I Solved Them

- Getting generic answers: Started using the brief template. Once I loaded 3–5 pieces of relevant context, answers became specific and usable.
- AI agreeing with everything I implied: Shifted from leading questions to neutral comparison requests. Instead of asking which is better, I asked for pros and cons of each.
- Suspecting the AI was guessing numbers instead of calculating: Added "write and run code, show me the code" to any data-related prompt. This forced actual computation.
- Not knowing if the first answer was good enough: Adopted the Brainstorm Iterate loop. I now plan for 3 rounds minimum options, feedback, and expand.
- Trusting one model too much: Started using two different AI tools for important work. Cross-checking caught errors and built confidence.

What Improved Final Outputs the Most

- Loading context upfront : Five lines of good context beat five paragraphs of clever wording.
- Iterating instead of accepting: The biggest quality jump came from rounds 2 and 3 of feedback.
- Forcing code execution for data: Eliminated silent guessing now all numerical work became verifiable.
- Using neutral framing : Got honest, balanced analysis instead of agreeable flattery.
- Cross-model verification : Caught blind spots a single model would have missed.